Strata.io vs. an AI Gateway

Detailed product comparison

Scope: This document compares Strata Identity's AI Identity Gateway at strata.io – https://www.strata.io/ with a model-traffic AI gateway, using Cloudflare AI Gateway as a concrete example. "AI gateway" is a broad category, and individual products differ, so the Cloudflare column illustrates a common model gateway rather than representing every vendor.

Executive summary

Strata's AI Identity Gateway and a model gateway both sit between applications and other systems, but they govern different parts of an AI workflow:

- Strata governs agent identity and access to tools, APIs, and MCP resources. It focuses on authentication, authorization, delegated credentials, and auditing agent actions.
- A model gateway governs calls to AI models. It commonly focuses on routing, provider fallbacks, rate limits, caching, spend controls, and request or response visibility.

They can be complementary. A model gateway can control which model receives an inference request, while Strata can control which business systems an agent may access and what it may do there. Strata describes its product as a runtime identity and policy enforcement point for agents and MCP environments. Strata announcement – https://www.strata.io/resources/news/strata-identity-introduces-ai-identity-gateway-and-validation-sandbox/ Cloudflare describes its gateway in terms of model request visibility, caching, rate limits, retries, and fallbacks. Cloudflare AI Gateway overview – https://developers.cloudflare.com/ai-gateway/

What each product is for

Strata Identity Gateway

Strata positions its AI Identity Gateway as part of Maverics for Agentic Identity, built on its broader identity orchestration platform. The gateway mediates an agent's access to upstream services, particularly MCP servers and APIs.

Strata documents a flow in which an agent presents a user access token to the gateway. The gateway can check policy for the requested tool, exchange the incoming token for a short-lived, narrowly scoped token, and forward the request to the upstream service. Its tutorial covers OIDC federation, inbound policy, token exchange, and token minting policy. Strata agentic identity flow – https://legacy-docs.strata.io/docs/build-

an-agentic-identity-flow

The key question Strata is designed to help answer is:

"Is this agent, acting for this user or workload, allowed to invoke this tool or reach this resource right now?"

Model-traffic AI gateway

A model gateway sits on the inference path between an application and one or more model providers. Cloudflare's AI Gateway, for example, supports analytics and logs, caching, rate limiting, request retries, and model fallbacks. Its dynamic routing can select models based on request conditions and metadata. Cloudflare AI Gateway overview – https://developers.cloudflare.com/ai-gateway/ Cloudflare dynamic routing – https://developers.cloudflare.com/ai-gateway/features/dynamic-routing/

The key question this type of gateway is designed to help answer is:

"Where should this model request go, and how should its cost, reliability, and usage be managed?"

Side-by-side comparison

Comparison at a glance

Primary control point

STRATA Agent-to-tool, agent-to-API, or agent-to-MCP access

MODEL GATEWAY Application-to-model inference requests

Core purpose

STRATA Enforce identity-aware authorization for agent actions

MODEL GATEWAY Manage model-provider traffic and operations

Main protected resource

STRATA Enterprise APIs, MCP tools, and downstream systems

MODEL GATEWAY Model endpoints and inference services

Identity model

STRATA Human and agent identity, often tied to an identity provider and user attributes

MODEL GATEWAY Usually application, project, user, or team metadata; identity depth varies by product

Authorization

STRATA Fine-grained policies can consider identity, tool, parameters, and context

MODEL GATEWAY Often access keys, quotas, and routing rules; some products also offer content or data guardrails

Credentials

STRATA Can exchange an inbound token for a short-lived token scoped to an upstream resource

MODEL GATEWAY Typically manages provider credentials or routes through a gateway credential

Routing

STRATA Routes agent tool calls to MCP servers or APIs

MODEL GATEWAY Routes model requests to providers, models, or fallback models

Reliability features

STRATA Identity continuity and controlled access are central to Strata's broader platform; gateway-specific model failover is not its primary role

MODEL GATEWAY Retries, provider/model fallbacks, and routing are common features

Cost controls

STRATA Not the central purpose of the identity gateway

MODEL GATEWAY Often includes request/token analytics, budgets, rate limits, or cost controls

Audit and observability

STRATA Records agent requests, responses, and token exchanges for access auditing

MODEL GATEWAY Records model request outcomes, errors, usage, and cost; retention and data handling vary

Deployment emphasis

STRATA Strata says the gateway runs in the customer's environment, including disconnected or air-gapped environments

MODEL GATEWAY Deployment varies: cloud service, self-hosted proxy, edge service, or a combination

Example integration

STRATA Agent → Strata → MCP server/API

MODEL GATEWAY Application → AI gateway → model provider

Cloudflare's feature documentation gives examples of model-gateway features including cost budgets, rate limiting, dynamic routing, guardrails, DLP, and logging. Which of those are included depends on the chosen gateway and plan. Cloudflare feature list – https://developers.cloudflare.com/ai-gateway/features/

Example architecture: using both

An agent may need a model to reason and a business system to act. Those are two separate access paths:

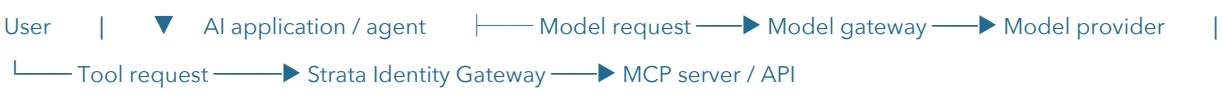

In this arrangement:

The **model gateway** manages inference traffic: for example, selecting a model, enforcing a usage limit, logging an error, or falling back to another provider.

Strata manages the agent's downstream access: for example, validating the agent's token, checking whether the user may invoke a tool, or obtaining a token with only the scopes needed for that call.

This separation matters because permission to call a model does not automatically establish permission to use a business tool. Likewise, permission to call a business tool does not determine which model should answer an inference request.

Where capabilities may overlap

The products can both be described as gateways, and both can contribute to governance and visibility. That does not make their controls interchangeable.

- Logging: Strata focuses on agent actions and identity/token events; model gateways commonly focus on inference requests and provider outcomes.
- Policy: Strata's documented policy path governs access to tools and upstream resources. A model gateway may apply different policies to routing, budgets, safety, or data handling.
- Authentication: Both products may authenticate traffic, but the credentials and trust relationships being checked can differ.
- Security controls: A model gateway with DLP or content moderation may inspect prompts and responses. Strata's documented use case is access control and delegated authorization for agent actions.

For example, prompt filtering does not itself authorize an agent to read a customer record. A least-privilege token for a customer API does not itself prevent sensitive text from being sent to an external model provider.

When Strata is the more relevant fit

Strata is worth evaluating when the core problem is:

- Agents access internal APIs or MCP tools using broad, static credentials.
- You need to make authorization decisions using the user or agent's identity and attributes.
- You want to exchange a broad incoming token for a short-lived, narrowly scoped upstream token.
- You need a policy and audit point for agent tool invocations.

- Your deployment requirements include running the enforcement gateway inside your own environment, including disconnected environments.

Strata's published tutorial describes inbound authorization with OPA policies and outbound token exchange for upstream resources. Strata agentic identity flow — https://legacy-docs.strata.io/docs/build-an-agentic-identity-flow

When a model gateway is the more relevant fit

A model gateway is worth evaluating when the core problem is:

- Applications call several model providers or models.
- You need provider or model routing and fallback behavior.
- You want a central place to inspect model usage, errors, or cost.
- You need rate limits, budgets, caching, retries, or request-level controls.
- You want provider-independent handling of model calls.

Cloudflare documents these types of features for its AI Gateway, including provider/model routing, analytics, caching, rate limits, retries, and fallbacks. Cloudflare overview — https://developers.cloudflare.com/ai-gateway/ Cloudflare features — https://developers.cloudflare.com/ai-gateway/features/

Evaluation questions

Questions for Strata

- Which identity providers and agent identity sources are supported in the intended environment?
- How are policies authored, reviewed, versioned, and rolled back?
- Which MCP transports and API integration patterns are supported?
- How are tool arguments incorporated into authorization decisions?
- How are exchanged tokens scoped, expired, and revoked?
- What audit events are captured, and how are they exported to the organization's SIEM?
- Which deployment, availability, and disconnected-operation requirements are supported by the purchased configuration?

Questions for a model gateway

- Which providers, models, and API formats are supported?
- Does routing support required rules, fallbacks, and streaming behavior?
- How are provider credentials stored and rotated?

- Can budgets and rate limits be scoped to teams, projects, or end users?
- Are prompts and completions retained, and can logging be disabled or redacted?
- Which DLP, moderation, and data residency controls are included?
- What happens to requests during gateway or provider outages?

Practical conclusion

Strata is an identity and authorization gateway for agent actions. A model gateway is a traffic-management layer for model calls. If your main risk is an agent accessing tools or data under excessive permissions, focus on Strata and the surrounding identity design. If your main operational issue is managing model providers, costs, and inference reliability, focus on a model gateway. If agents both call models and take actions through enterprise tools, evaluate whether you need both control points and verify how their logs and policies fit together.

This comparison describes publicly documented product positioning and capabilities. Exact features, deployment options, and commercial availability should be confirmed with the vendors for your use case.